

Automatic differentiation/backpropagation

loss.backward() →

param.grad ← what grad? why?
1.bw?

Computational Differentiation

① Numerical Differentiation

very good test

Disadvantages: *cute to verify the correctness*

a) Approximate *of your implementation*

b) Computationally expensive:

How many times do you have to call

$f(\underline{x})$ to compute $\frac{\partial f(\underline{x})}{\partial \underline{x}}$ if $\underline{x} \in \mathbb{R}^n$?

$n+1$ times

$$f(\underline{x})$$

$$\frac{\partial f(\underline{x})}{\partial \underline{x}} = \left[\frac{\partial f(\underline{x})}{\partial x_1}, \dots, \frac{\partial f(\underline{x})}{\partial x_n} \right]$$

$$\frac{\partial f(\underline{x})}{\partial x_i} =$$

$$\frac{f(\underline{x} + \underline{\epsilon}_i \underline{\epsilon}) - f(\underline{x})}{\underline{\epsilon}}$$

$$\underline{\epsilon}_i = \begin{cases} 0 \\ \vdots \\ i \\ \vdots \\ 0 \end{cases}$$

1 at ith place everywhere

② Symbolic Differentiation (SD)

e.g. sympy / Matlab

formula
of a function



Disadvantages

- a) you need a new language to communicate your formulas
- b) You get formula as output. Efficiency is not considered

formula
for the derivative

→ to implement
the derivative
comp. complex?

③ Automatic differentiation (AD)

→ Operator overloading to implement (AD)

C++ / Python

$a + b$
mt mt

$a = \text{Tensor}()$
 $b = \text{Tensor}()$

$a = \text{Person}()$

$b = \text{Person}()$

$a + b$

$a @ b$

$\text{Person}.___add__$

$\text{Person}.___matmul__$

```

class Tensor():
    def __init__(self):
        self.value = 0
        self.grad = 0

```

a = tensor(10)
 $\circ = \text{Tensor}(10)$

$$c = a + b$$

$$c.value = a.value + b.value$$

$$c.grad = \underline{\hspace{1cm}}$$

$a = \text{Symbol('a')}$

$b = \text{Symbol('b')}$

$$c = a + b$$

$$c.expression = \underline{\text{"a+b"}}$$

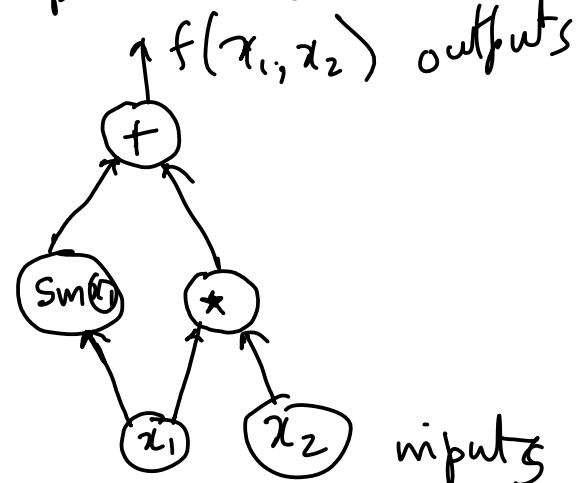
AD

- ① Forward-mode AD (Forward accumulative)
- ② Reverse-mode AD (Backpropagation) in NN

Mathematical expressions can be represented as a Directed Acyclic graph (computation graph)

Example

$$f(x_1, x_2) = x_1 x_2 + \sin(x_1)$$

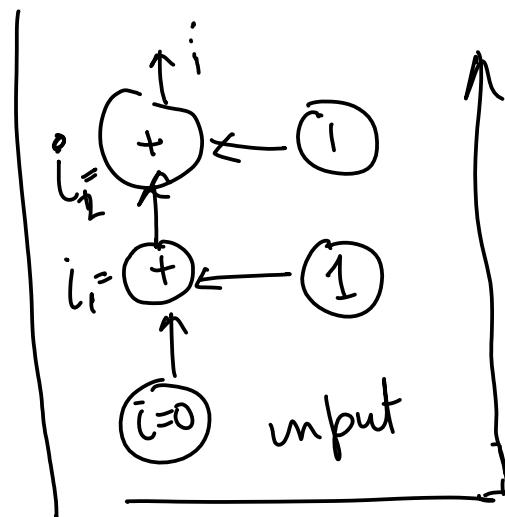


outputs

```

i = Tensor(0)
for i in range(10)
    i = i + 1

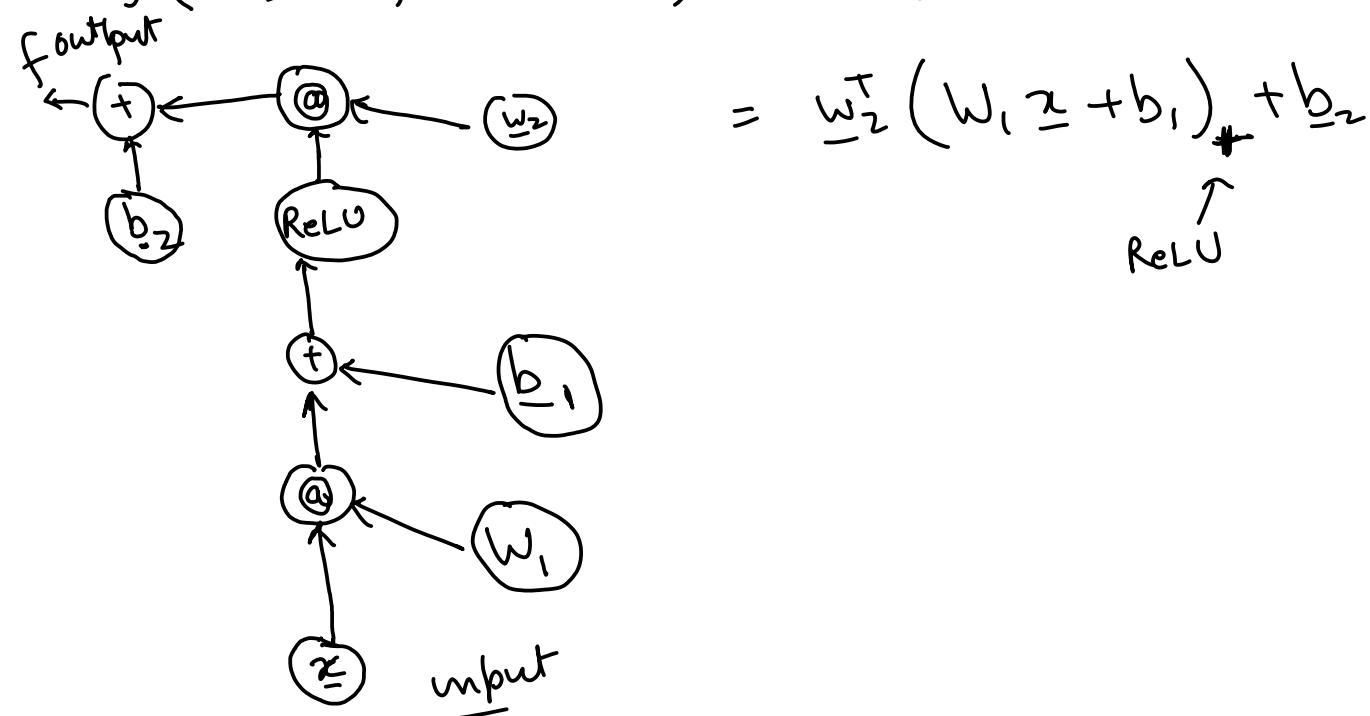
```



Internally
Pytorch
will
keep a
track of
this graph

Computation Graph of MLP (2-layer, ReLU)

$$f(\underline{x}; \underline{w}_1, \underline{w}_2, \underline{b}_1, \underline{b}_2) = \underline{w}_2^T \text{ReLU}(\underline{w}_1 \underline{x} + \underline{b}_1) + \underline{b}_2$$



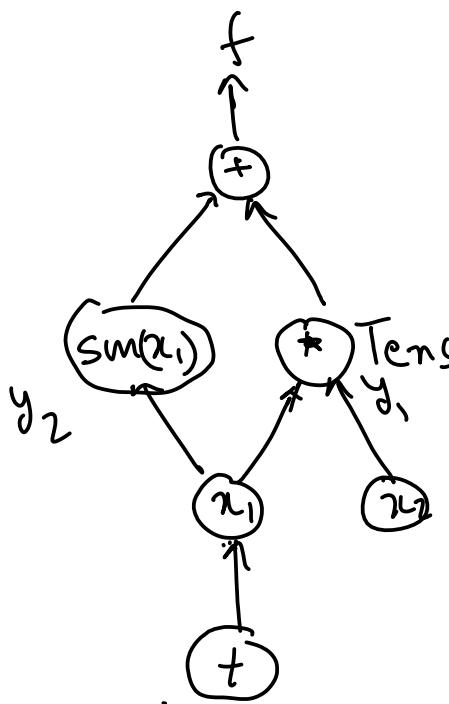
(1) Forward-mode AD

Use operator overloading to implement AD

```
class Tensor():
    def __init__(self):
        self.value = None
        self.grad = None
```

a = Tensor(10)
b = Tensor(20)

$$c = a + b$$



class Tensor

```
def __add__(self, b):
    c = Tensor()
    c.value = self.value + b.value
    c.grad = self.grad + b.grad
    return c
```

value
grad = $\frac{\partial f}{\partial y_1}$. $\frac{\partial y_1}{\partial x_1} = ?$

Reverse mode

Forward mode

In forward mode

$$x_1.grad = \frac{\partial x_1}{\partial t} = 1$$

$$x_2.grad = \frac{\partial x_2}{\partial t} = 0$$

$$y_1 = x_1 * x_2 \Rightarrow y_1.grad = \frac{\partial (x_1 * x_2)}{\partial t} = x_1 \frac{\partial x_2}{\partial t} + x_2 \frac{\partial x_1}{\partial t}$$

$$\Rightarrow y_1 \cdot \text{grad} = x_1 \cdot (x_2 \cdot \text{grad}) + x_2 \cdot (x_1 \cdot \text{grad})$$

$$y_2 = \sin(x_1) \rightarrow \frac{\partial y_2}{\partial t} = y_2 \cdot \text{grad} = \cos(x_1) \frac{\partial x_1}{\partial t} = \cos(x_1) (x_1 \cdot \text{grad})$$

$$f = y_1 + y_2 \rightarrow \frac{\partial f}{\partial t} = f \cdot \text{grad} = \frac{\partial y_1}{\partial t} + \frac{\partial y_2}{\partial t}$$

$$= y_1 \cdot \text{grad} + y_2 \cdot \text{grad}$$


Accumulation direction

$$\frac{\partial f}{\partial t} = \left\{ \frac{\partial}{\partial y_1} (y_1 + y_2) \left[\frac{\partial (x_1 \cdot x_2)}{\partial x_1}, \frac{\partial x_1}{\partial t} \right] \right\}$$

This accumulation happens first

$$f = h(g(k(l(x))))$$

$$\frac{\partial f}{\partial x} = \left(\frac{\partial h}{\partial g} \left(\frac{\partial g}{\partial k} \left(\frac{\partial k}{\partial l} \frac{\partial l}{\partial x} \right) \right) \right)$$




- ① In forward mode differentiation, the accumulation of chain rule derivatives happens from the input side to the output side (forward direction).
- ② In forward mode differentiation, you have to re-propagate the derivatives for each input.

$$f_1(x_1, x_2) = x_1 \cdot x_2 + \sin(x_1)$$

$$f_2(x_1, x_2) = x_1 \cdot x_2 + \cos(x_2)$$

$$\frac{\partial f_1}{\partial t}, \frac{\partial f_2}{\partial t}$$




- ③ we don't have to recompute derivatives for the common part of the graph.

② Reverse-mode AD


$$L(D; w) = \sum_D \|y - f(x, w)\|_2^2$$

scalar





- ① Reverse mode is more efficient than forward mode when # of output < # of inputs

Forward mode AD




Reverse mode AD



NIP-hard

chain
graph


$$\begin{aligned} f &= f = \cdot \\ h &= \begin{bmatrix} \text{sm}(x_1) \\ x_1 * x_2 \end{bmatrix} \\ x &= \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \end{aligned}$$



$$\frac{\partial f}{\partial x} = \underbrace{\frac{\partial f}{\partial h}}_{\text{vector}} \underbrace{\frac{\partial h}{\partial x}}_{\text{Matrix}} \quad \text{Jacobian}$$

$$f = f(g(h(k(z))))$$

$$\frac{\partial f}{\partial x_i} = \frac{\partial f}{\partial g} \cdot \frac{\partial g}{\partial h} \circ \frac{\partial h}{\partial k} \circ \frac{\partial k}{\partial x_i}$$



$$\frac{\partial f}{\partial x_i} = \left(\frac{\partial f}{\partial g} \left(\frac{\partial g}{\partial h} \left(\frac{\partial h}{\partial k} \cdot \frac{\partial k}{\partial x_i} \right) \right) \right)$$

$$\frac{\partial f}{\partial x_i} = \left(\frac{\partial f}{\partial g} \left(\frac{\partial g}{\partial h} \left(\frac{\partial h}{\partial k} \cdot \frac{\partial k}{\partial x_i} \right) \right) \cdot \frac{\partial h}{\partial k} \cdot \frac{\partial k}{\partial x_i} \right)$$

$\mathcal{O}(\text{matrix size})$

output < # inputs

$$\mathcal{O}_{FD}() > \mathcal{O}_{RD}()$$

Before you turn this problem in, make sure everything runs as expected. First, **restart the kernel** (in the menubar, select Kernel→Restart) and then **run all cells** (in the menubar, select Cell→Run All).

Make sure you fill in any place that says YOUR CODE HERE or "YOUR ANSWER HERE", as well as your name and collaborators below:

```
In [ ]: NAME = ""  
COLLABORATORS = ""
```


Differentiation options

1. Numerical differentiation
2. Symbolic differentiation
3. Automatic differentiation
 - A. Forward mode differentiation
 - B. Reverse mode differentiation

1. Numerical differentiation

2. Symbolic differentiation

3. Automatic differentiation



3.A Forward mode

Example:

$$z = f(x_1, x_2) = x_1 x_2 + \sin(x_1)$$

3.B Reverse mode

Example:

$$z = f(x_1, x_2) = x_1x_2 + \sin(x_1)$$

```
In [ ]: import numpy as np
class ForwardDiff:
    def __init__(self, value, grad=None):
        self.value = value
        self.grad = np.zeros_like(value) if grad is None else grad

    def __add__(self, other):
        cls = type(self)
        other = other if isinstance(other, cls) else cls(other)
        out = cls(self.value + other.value,
                  self.grad + other.grad)
        return out
    __radd__ = __add__

    def __repr__(self):
        return f"{self.__class__.__name__}(data={self.value}, grad={self.grad})"

x = ForwardDiff(2, 1)
y = ForwardDiff(3, 0)

f = x + y
f
```

```
In [ ]: oldFD = ForwardDiff # Bad practice: do not do it
class ForwardDiff(oldFD):
    def __mul__(self, other):
        cls = type(self)
        other = other if isinstance(other, cls) else cls(other)
        out = cls(self.value * other.value,
                  other.value * self.grad +
                  self.value * other.grad)
        return out

    __rmul__ = __mul__

x = ForwardDiff(2, 0)
y = ForwardDiff(3, 1)

f1 = x * y
f2 = 2*x + 3*y + x*y
f1, f2
```

```
In [ ]: oldFD = ForwardDiff # Bad practice: do not do it
class ForwardDiff(oldFD):
    def log(self):
        cls = type(self)
        return cls(np.log(self.value),
                  1/self.value * self.grad)
    def exp(self):
        cls = type(self)
```

```

        out_val = np.exp(self.value)
        return cls(out_val,
                   out_val * self.grad)

    def sin(self):
        cls = type(self)
        return cls(np.sin(self.value),
                   np.cos(self.value) * self.grad)

    def cos(self):
        cls = type(self)
        return cls(np.cos(self.value),
                   -np.sin(self.value) * self.grad)

    def __pow__(self, other):
        cls = type(self)
        other = other if isinstance(other, cls) else cls(other)
        return (self.log() * other).exp()

    def __neg__(self): # -self
        return self * -1

    def __sub__(self, other): # self - other
        return self + (-other)

    def __truediv__(self, other): # self / other
        return self * other**-1

    def __rtruediv__(self, other): # other / self
        return other * self**-1

```



```

x = ForwardDiff(2, 1)
y = ForwardDiff(3, 0)

f = x**y
f

```

```

In [ ]: import numpy as np
def add_vjp(a, b, grad):
    return grad, grad

def no_parents_vjp(grad):
    return (grad,)

class ReverseDiff:
    def __init__(self, value, parents=(), op='', vjp=no_parents_vjp):
        self.value = value
        self.parents = parents
        self.op = op
        self.vjp = vjp
        self.grad = None

    def backward(self, grad):

```

```

        self.grad = grad
        op_args = [p.value for p in self.parents]
        grads = self.vjp(*op_args, grad)
        for g, p in zip(grads, self.parents):
            p.backward(g)

    def __add__(self, other):
        cls = type(self)
        other = other if isinstance(other, cls) else cls(other)
        out = cls(self.value + other.value,
                  parents=(self, other),
                  op='+',
                  vjp=add_vjp)
        return out

    __radd__ = __add__

    def __repr__(self):
        cls = type(self)
        return f"{cls.__name__}(value={self.value}, parents={self.parents},"

x = ReverseDiff(2)
y = ReverseDiff(3)

f = x + y + 3
f.backward(1)
f
x.grad, y.grad

```

In []: oldRD = ReverseDiff # Bad practice: do not do it

```

def mul_vjp(a, b, grad):
    return grad * b, grad * a

class ReverseDiff(oldRD):
    def __mul__(self, other):
        cls = type(self)
        other = other if isinstance(other, cls) else cls(other)
        out = cls(self.value * other.value,
                  parents=(self, other),
                  op='*',
                  vjp=mul_vjp)
        return out

    __rmul__ = __mul__

x = ReverseDiff(2)
y = ReverseDiff(3)

f1 = 5*x + 7*y
f1.backward(1)
x.grad, y.grad

```

In []: f2 = x*y
f2.backward(1)

```
x.grad, y.grad
```